a first order buys a bit per token; the control pays for it

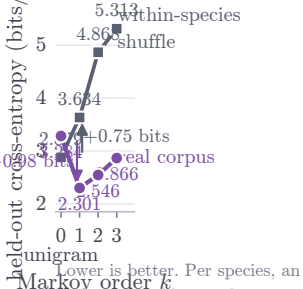

Lower is better. Per species, an 80/20 split by bout, a k-order backoff model fitted on the first half, every held-out token scored. An unbiased estimator paying for parameters it does not use. The clean form of the claim that the tokens are units. Production corpus: 34 species, 327 164 tokens.

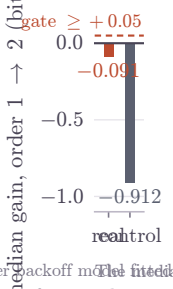

The median species loses 0.091 bits per token going to second order, at three tokens of context (−0.091 / −0.116 / −0.142). The corpus cannot estimate the contexts a second order asks for.

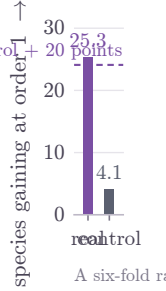

A six-fold rate over the shuffle, clearing the 20-point margin. So: there is no second-order structure this corpus can support, and a quarter of species hint at one. 217 of 2 421 species can be asked at all.