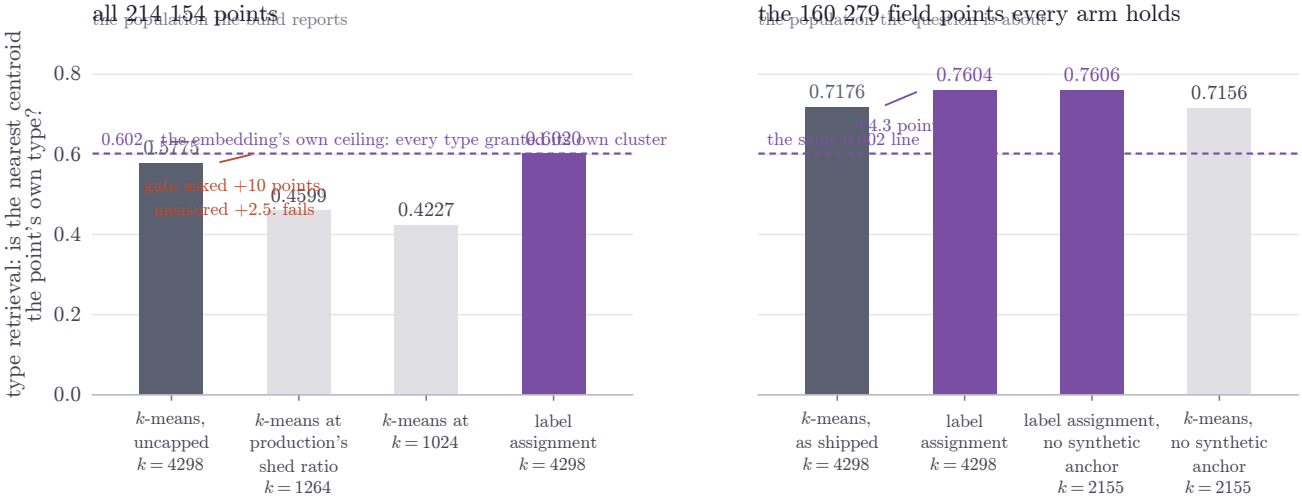

One fixed matrix dumped from a 100-species debug build: 214 154 points, 2 143 types, 160 279 of the points from field recordings; every arm scored through the shipped clustering, seeding and purity functions. Changing the query population moves the same arm by 14 to 16 points, which is larger than any difference between arms - and an arm that drops the synthetic copies drops the hard queries, so on all points it flatters itself by 18. Purity and type coverage are identities under label assignment, not measurements; type retrieval is the statistic that survives because nothing forces it. Debug-scale.