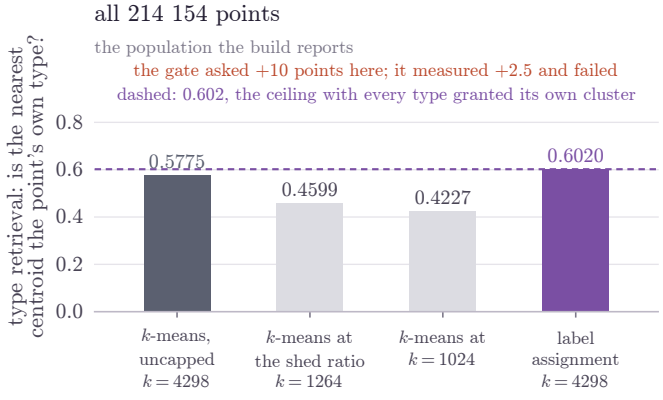

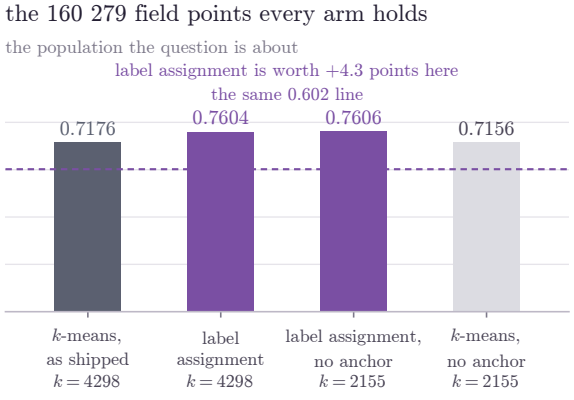

The question is the one the browser asks of live audio: is the nearest centroid by cosine the point's own type? One fixed matrix dumped from a 100-species debug build - 214 154 points, 2 143 types, 160 279 of them from field recordings - with every arm scored through the shipped clustering, seeding and purity functions. On a corpus that fits under the cluster cap, seeded k-means is nearly as good as knowing the answer, which is why the gate fails and why the failure is the useful part. Changing the query population moves the same arm by 14 to 16 points, more than any arm moves it, and an arm that drops the synthetic copies drops the hard queries, so on all points it flatters itself by 18. Purity and coverage are identities under label assignment, not measurements. Debug-scale.