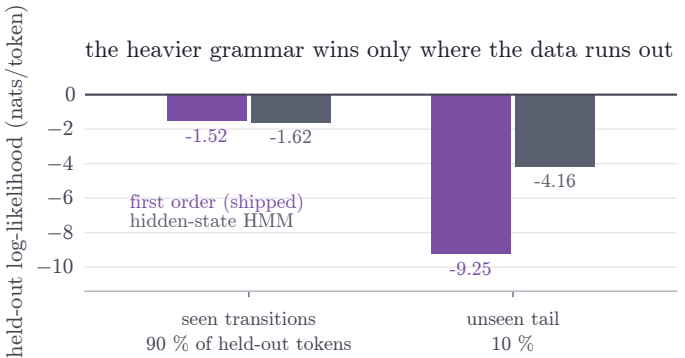


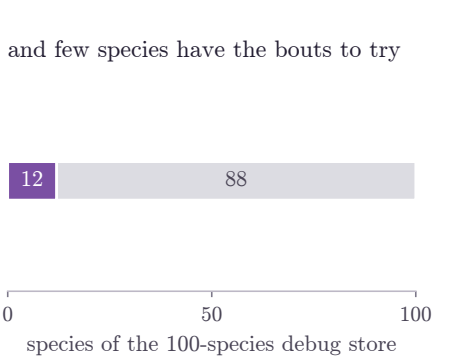


Left: first order is the better model on the 90 % of held-out tokens whose transition was seen in training, and the hidden-state model's entire advantage is the unseen tail. That is a property of smoothing shape, not of recovered hidden structure, so the smoothing was repaired instead: a Kneser-Ney continuation backoff gains 1.90 and 2.32 nats per token on unseen-context transitions across two stores. Right: twelve species reach 20 bouts after run-collapse, which is the syntax threshold. Counted before the depth harvest; the production re-count is 217 of 2 421 species. Both panels are debug-scale.