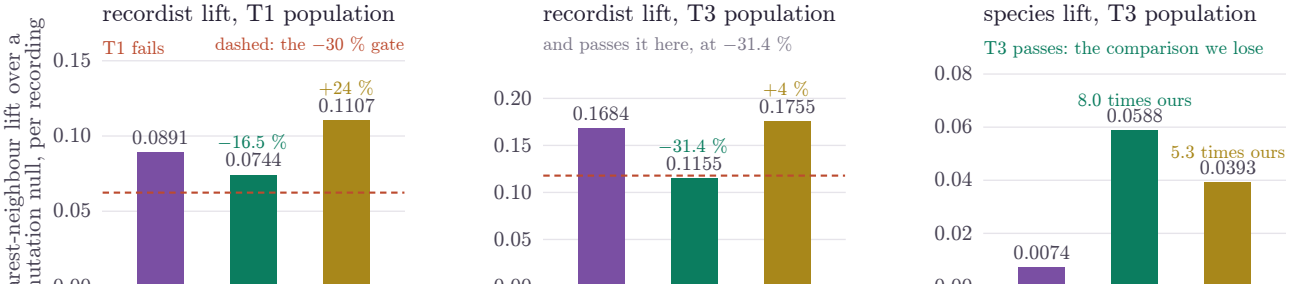

Read-only tests on a fixed store, one estimator, no training and no new data. The representation is per recording in every row, so same-recording leakage is impossible by construction - which also means a species result here says an external space is the better per-recording descriptor, and does not say Perch beats our critic syllable for syllable, a comparison that has never been run. T1's verdict rests on a coverage tie-break: it passes at −31.4 % on the T3 population and fails at −16.5 % on its own, whose recordist estimate covers 0.907 of rows against 0.329, so the instrument built for the question wins. The finding that holds in both draws is the one to carry: the recordist lift is twice the species lift in Perch's space, 4.5 times in BirdNET's and 23 times in ours. Every space tested is confounded in the same direction and they differ only in degree.