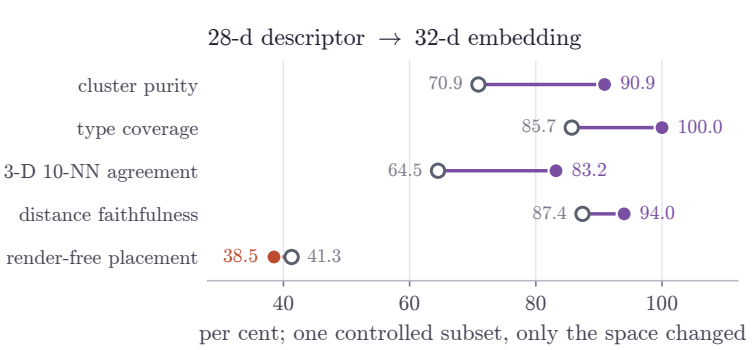

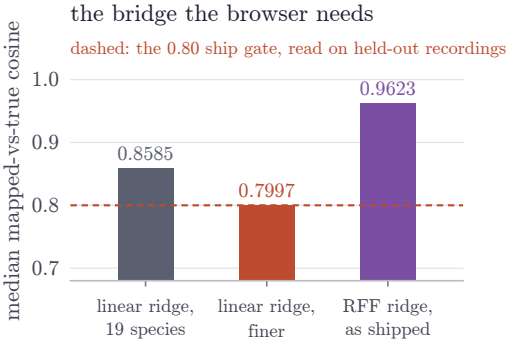

Left: 625 field syllables, three species, identical points, one seed. Hollow marker, the 28-d descriptor; filled, the 32-d embedding. Clay is the counterweight, because render-free placement moves the wrong way and the recordist confound lives in the embedding rather than in the descriptor: the embedding is the better similarity space and the worse-behaved one. Right: the browser computes only the 28-d descriptor, so a kernel ridge over random Fourier features ($D = 2048$, $\sigma = 4$, $\lambda = 1$) maps it into the 32-d space the map lives in. It is gated on held-out recordings, never on held-out syllables, because two syllables of one bird in one take are nearly the same pair twice; it ships as plain numbers, and no weights reach the browser. Both panels are debug-scale.